



Project By : PRASAD JADHAV

```
In [33]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from imblearn.over_sampling import SMOTE

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn import svm

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: dataset = pd.read_csv('diabetes.csv')
dataset.shape
```

```
Out[3]: (768, 9)
```

```
In [4]: dataset.head()
```

```
Out[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [5]: dataset.tail()
```

```
Out[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	0
767	1	93	70	31	0	30.4	0.315	23	0

```
In [6]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null    int64
1   Glucose                768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
4   Insulin                768 non-null    int64
5   BMI                    768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                    768 non-null    int64
8   Outcome                768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [7]: dataset.describe()
```

Out[7]:		Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
	<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
	<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.332424
	<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.332671
	<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
	<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.000000
	<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.000000
	<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.000000
	<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	0.671000


In [8]: `dataset.corr()`

Out[8]:		Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
	<b>Pregnancies</b>	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	0.332424
	<b>Glucose</b>	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.332671
	<b>BloodPressure</b>	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.332671
	<b>SkinThickness</b>	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.332671
	<b>Insulin</b>	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.332671
	<b>BMI</b>	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.332671
	<b>DiabetesPedigreeFunction</b>	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000
	<b>Age</b>	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.332671
	<b>Outcome</b>	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	0.332671


In [9]: `dataset.cov()`

Out[9]:		Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
	<b>Pregnancies</b>	11.354056	13.947131	9.214538	-4.390041	-28.555231	0.017683	0.332424
	<b>Glucose</b>	13.947131	1022.248314	94.430956	29.239183	1220.935799	0.221071	0.332671
	<b>BloodPressure</b>	9.214538	94.430956	374.647271	64.029396	198.378412	0.281805	0.332671
	<b>SkinThickness</b>	-4.390041	29.239183	64.029396	254.473245	802.979941	0.392573	0.332671
	<b>Insulin</b>	-28.555231	1220.935799	198.378412	802.979941	13281.180078	0.197859	0.332671
	<b>BMI</b>	0.469774	55.726987	43.004695	49.373869	179.775172	1.000000	0.332671
	<b>DiabetesPedigreeFunction</b>	-0.037426	1.454875	0.264638	0.972136	7.066681	0.140647	1.000000
	<b>Age</b>	21.570620	99.082805	54.523453	-21.381023	-57.143290	0.036242	0.332671
	<b>Outcome</b>	0.356618	7.115079	0.600697	0.568747	7.175671	0.292695	0.332671

In [10]: `sns.distplot(dataset['Outcome'])`  
`plt.show()`



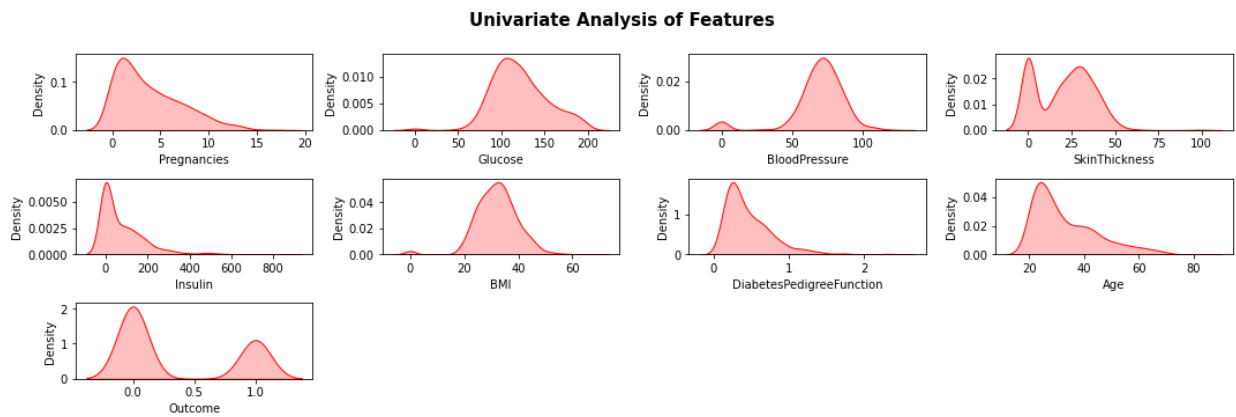
```
In [11]: sns.scatterplot(dataset['Age'],dataset['Outcome'])
plt.show()
```



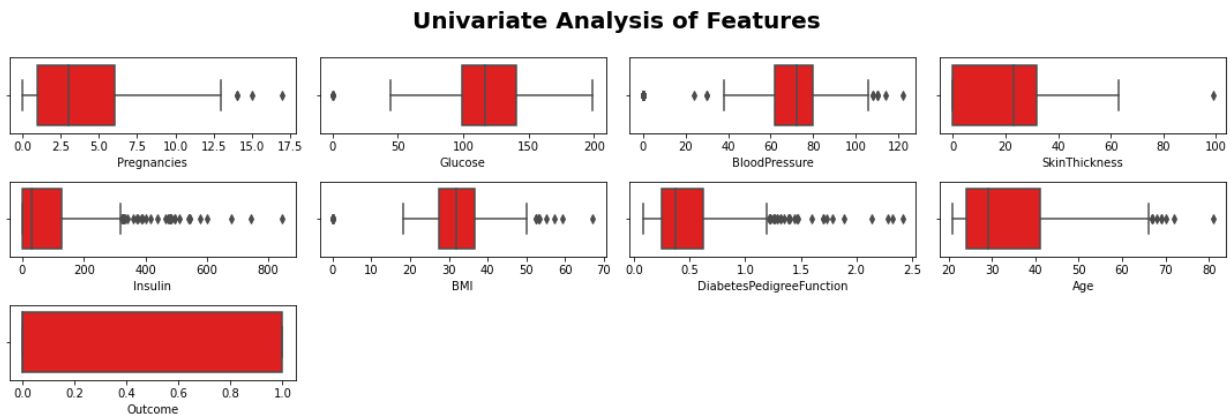
```
In [12]: features = [feature for feature in dataset.columns]
```

```
In [13]: plt.figure(figsize=(15,15))
plt.suptitle('Univariate Analysis of Features',fontweight='bold',fontsize=15,y=1.05)

for i in range(0,len(features)):
    plt.subplot(10,4,i+1)
    sns.kdeplot(x=dataset[features[i]],shade=True,color='red')
plt.tight_layout()
```



```
In [14]: plt.figure(figsize = (15,15))
plt.suptitle('Univariate Analysis of Features',fontweight='bold',fontsize=20,y=
for i in range(0,len(features)):
    plt.subplot(10,4,i+1)
    sns.boxplot(data=dataset,x=features[i],color='red')
    plt.xlabel(features[i])
    plt.tight_layout()
```



```
In [15]: dataset['Outcome'].value_counts()
```

```
Out[15]: 0    500
         1    268
         Name: Outcome, dtype: int64
```

```
In [16]: X = dataset.drop('Outcome',axis=1)
         y = dataset['Outcome']
```

```
In [19]: X_res,y_res = SMOTE().fit_resample(X,y)
```

```
In [20]: X_train,X_test,y_train,y_test = train_test_split(X_res,y_res,test_size=0.20,ra
```

```
In [21]: st = StandardScaler()
         X_train = st.fit_transform(X_train)
         X_test = st.fit_transform(X_test)
```

```
In [22]: model_df = {}

def model_val(model,X,y):
    X_train,X_test,y_train,y_test = train_test_split(X_res,y_res,test_size=0.20
```

```

model.fit(X_train,y_train)
y_pred = model.predict(X_test)
print(f'{model} Accuracy is {accuracy_score(y_test,y_pred)}')

score = cross_val_score(model,X,y,cv=5,n_jobs=-1)
print(f'{model} Average cross val score is {np.mean(score)}')
model_df[model] = round(np.mean(score)*100,2)

```

```

In [23]: model = LogisticRegression()
model_val(model,X,y)

```

```

LogisticRegression() Accuracy is 0.71
LogisticRegression() Average cross val score is 0.7695696460402341

```

```

In [24]: model = DecisionTreeClassifier()
model_val(model,X,y)

```

```

DecisionTreeClassifier() Accuracy is 0.7
DecisionTreeClassifier() Average cross val score is 0.7123673711909005

```

```

In [25]: model = RandomForestClassifier()
model_val(model,X,y)

```

```

RandomForestClassifier() Accuracy is 0.78
RandomForestClassifier() Average cross val score is 0.7683133859604447

```

```

In [26]: model = GradientBoostingClassifier()
model_val(model,X,y)

```

```

GradientBoostingClassifier() Accuracy is 0.77
GradientBoostingClassifier() Average cross val score is 0.7591715474068416

```

```

In [27]: model_df

```

```

Out[27]: {LogisticRegression(): 76.96,
DecisionTreeClassifier(): 71.24,
RandomForestClassifier(): 76.83,
GradientBoostingClassifier(): 75.92}

```

```

In [28]: C = 1.0
svm = svm.SVC(kernel='linear',C=C)
svm.fit(X_train,y_train)

```

```

Out[28]: SVC
SVC(kernel='linear')

```

```

In [29]: y_pred = svm.predict(X_test)

```

```

In [30]: accuracy_score(y_test,y_pred)

```

```

Out[30]: 0.705

```

```
In [31]: parameters = [
    {'kernel': ['linear'], 'C': [1, 10]},
    {'kernel': ['poly'], 'C': [1, 10]},
    {'kernel': ['rbf'], 'gamma': [1e-3, 1e-4], 'C': [1, 10]}
]
```

```
In [34]: gr = GridSearchCV(svm.SVC(), parameters, n_jobs=-1)
gr.fit(X_train, y_train)
```

```
Out[34]: ▸ GridSearchCV
          ▸ estimator: SVC
            ▸ SVC
```

```
In [36]: y_pred_ = gr.predict(X_test)
```

```
In [37]: accuracy_score(y_test, y_pred_)
```

```
Out[37]: 0.705
```

```
In [38]: from sklearn.model_selection import StratifiedKFold
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.calibration import CalibratedClassifierCV
from sklearn.neural_network import MLPClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.linear_model import SGDClassifier
from sklearn.svm import SVC
from xgboost import XGBClassifier
```

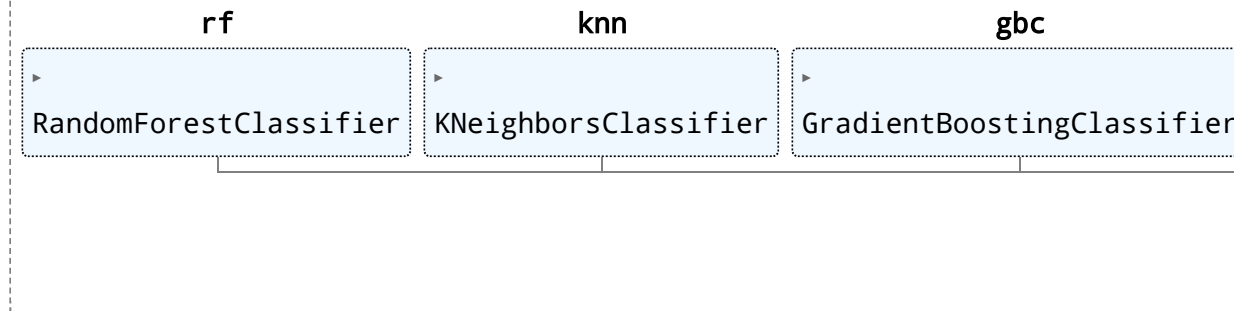
```
In [41]: estimators = [
    ('rf', RandomForestClassifier(n_estimators = 10, random_state = 42)),
    ('knn', KNeighborsClassifier(n_neighbors = 10)),
    ('gbc', GradientBoostingClassifier()),
    ('lr', LogisticRegression()),
    ('ccv', CalibratedClassifierCV()),
    ('mlp', MLPClassifier()),
    ('dt', DecisionTreeClassifier()),
    ('lda', LinearDiscriminantAnalysis()),
    ('gnb', GaussianNB()),
    ('adb', AdaBoostClassifier()),
    ('etc', ExtraTreesClassifier()),
    ('sgd', SGDClassifier()),
    ('svm', SVC()),
    ('xgb', XGBClassifier(n_estimators= 10, random_state = 42))
]
```

```
In [42]: from sklearn.ensemble import StackingClassifier
```

```
st = StackingClassifier(  
    estimators = estimators,  
    final_estimator = LogisticRegression(),  
    cv = 10,  
    stack_method='predict',  
    n_jobs=-1  
)
```

```
In [43]: st.fit(X_train,y_train)
```

```
Out[43]:
```



```
In [44]: y_pred = st.predict(X_test)
```

```
In [45]: accuracy_score(y_test,y_pred)
```

```
Out[45]: 0.785
```

```
In [46]: from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import BaggingClassifier  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.ensemble import ExtraTreesClassifier  
from sklearn.model_selection import GridSearchCV  
from xgboost import XGBClassifier
```

```
In [49]: dt_parameters = {  
    'criterion' : ['gini', 'entropy', 'log_loss'],  
    'splitter' : ['best', 'random'],  
    'max_depth' : range(1,10,1),  
    'min_samples_split' : range(2,10,2),  
    'min_samples_leaf' : range(1,5,1),  
    'max_features' : ['auto', 'sqrt', 'log2']  
}  
  
bg_parameters = {  
    'n_estimators' : [5, 10, 15],  
    'max_samples' : range(2, 10, 1),  
    'max_features' : range(2, 10, 3)  
}  
  
rf_parameters = {  
    'criterion' : ['gini', 'entropy', 'log_loss'],  
    'max_depth' : range(1, 10, 1),  
    'min_samples_split' : range(2, 10, 2),  
    'min_samples_leaf' : range(1, 10, 1),
```



```

}

et_parameters = {
    'n_estimators' : [10,20,30],
    'criterion' : ['gini', 'entropy', 'log_loss'],
    'max_depth' : range(2,10,1),
    'min_samples_split' : range(2,10,2),
    'min_samples_leaf' : range(1,5,1),
    'max_features' : ['sqrt', 'log2']
}

xgb_parameters = {
    'max_depth' : [2,4,6],
    'min_child_weight' : [2,4,6],
    'gamma' : [i/10 for i in range(4)]
}

```

```

In [50]: hyper_1 = GridSearchCV(estimator = DecisionTreeClassifier(), param_grid = dt_parameters)
hyper_2 = GridSearchCV(estimator = BaggingClassifier(), param_grid = bg_parameters)
hyper_3 = GridSearchCV(estimator = RandomForestClassifier(), param_grid = rf_parameters)
hyper_4 = GridSearchCV(estimator = ExtraTreesClassifier(), param_grid = et_parameters)
hyper_5 = GridSearchCV(estimator = XGBClassifier(), param_grid = xgb_parameters)

```

```

In [51]: hyper_1.fit(X_train,y_train)
hyper_2.fit(X_train,y_train)
hyper_3.fit(X_train,y_train)
hyper_4.fit(X_train,y_train)
hyper_5.fit(X_train,y_train)

```

Fitting 3 folds for each of 2592 candidates, totalling 7776 fits  
 Fitting 3 folds for each of 72 candidates, totalling 216 fits  
 Fitting 3 folds for each of 972 candidates, totalling 2916 fits  
 Fitting 3 folds for each of 2304 candidates, totalling 6912 fits  
 Fitting 3 folds for each of 36 candidates, totalling 108 fits

```

Out[51]:
└─ GridSearchCV
  └─ estimator: XGBClassifier
    └─ XGBClassifier

```

```

In [52]: h_pred_1 = hyper_1.predict(X_test)
h_pred_2 = hyper_2.predict(X_test)
h_pred_3 = hyper_3.predict(X_test)
h_pred_4 = hyper_4.predict(X_test)
h_pred_5 = hyper_5.predict(X_test)

```

```

In [53]: print(accuracy_score(y_test,h_pred_1))
print(accuracy_score(y_test,h_pred_2))
print(accuracy_score(y_test,h_pred_3))
print(accuracy_score(y_test,h_pred_4))
print(accuracy_score(y_test,h_pred_5))

```

0.705  
 0.67  
 0.795  
 0.795  
 0.78

```
In [54]: import pickle
import joblib
```

```
In [56]: pickle.dump(hyper_3,open('diabetes_prediction.pkl','wb'))
```

```
In [57]: model = pickle.load(open('diabetes_prediction.pkl','rb'))
```

```
In [58]: dataset.head()
```

```
Out[58]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [60]: new_df = pd.DataFrame({
    'Pregnancies':6,
    'Glucose':148,
    'BloodPressure':72,
    'SkinThickness':35,
    'Insulin':0,
    'BMI':33.6,
    'DiabetesPedigreeFunction':0.627,
    'Age':50,
    },index=[0])
```

```
In [61]: model.predict(new_df)
```

```
Out[61]: array([1], dtype=int64)
```

```
In [63]: p = hyper_3.predict(new_df)

prob = hyper_3.predict_proba(new_df)
if p == 1:
    print('Diabetes!')
    print(f'You will be Diabetes! with Probability of {prob[0][1]:.2f}')
else:
    print('Not-Diabetes!')
```

```
Diabetes!
You will be Diabetes! with Probability of 0.65
```